



WIKIPEDIA
The Free Encyclopedia

WIKIPEDIA

Mistral AI

Mistral AI SAS is a French artificial intelligence (AI) startup, headquartered in Paris. It specializes in open-weight large language models (LLMs).^{[2][3]}

Namesake

The company is named after the mistral, a powerful, cold wind in southern France.^[4]

History

Mistral AI was established in April 2023 by three French AI researchers, Arthur Mensch, Guillaume Lample and Timothée Lacroix.^[5]

Mensch, an expert in advanced AI systems, is a former employee of Google DeepMind; Lample and Lacroix, meanwhile, are large-scale AI models specialists who had worked for Meta Platforms.^[6]

The trio originally met during their studies at École Polytechnique.^[4]

Company operation

Philosophy

Mistral AI emphasizes openness and innovation in the AI field and positions itself as an alternative to proprietary models.^[7]

The company has gained prominence as an alternative to proprietary AI systems as it aims to "democratize" AI by focusing on open-source innovation.^[8]

Mistral AI SAS



Company type	Private
Industry	Artificial intelligence
Founded	28 April 2023
Founders	Arthur Mensch Guillaume Lample Timothée Lacroix
Headquarters	Paris, France
Key people	Arthur Mensch (CEO) Guillaume Lample (Chief Scientist) Timothée Lacroix (CTO)
Products	Mistral 7B Mixtral 8x7B Mistral Medium Mistral Large Mistral Large 2 (123B) Mixtral 8x22B Codestral 22B Codestral Mamba (7B) Mathstral (7B) Mistral NeMo 12B Mistral Embed
Number of employees	150 (2025) ^[1]
Website	<u>mistral.ai</u> (<u>https://mistral.ai/</u>)

Funding

In June 2023, the start-up carried out a first fundraising of €105 million (\$117 million) with investors including the American fund Lightspeed Venture Partners, Eric Schmidt, Xavier Niel and JCDecaux. The valuation is then estimated by the Financial Times at €240 million (\$267 million).

On 10 December 2023, Mistral AI announced that it had raised €385 million (\$428 million) as part of its second fundraising. This round of financing involves the Californian fund Andreessen Horowitz, BNP Paribas and the software publisher Salesforce.^[9]

In October 2023, Mistral AI raised €385 million.^[10]

By December 2023, it was valued at over \$2 billion.^{[11][12][13]}

On 16 April 2024, reporting revealed that Mistral was in talks to raise €500 million, a deal that would more than double its current valuation to at least €5 billion.^[14]

In June 2024, Mistral AI secured a €600 million (\$645 million) funding round, elevating its valuation to €5.8 billion (\$6.2 billion).^[15]


Led by venture capital firm General Catalyst,^[16] this round resulted in additional contributions from existing investors. The funds aim to support the company's expansion.

Based on valuation, the company is currently in fourth place in the global AI race and in first place outside the San Francisco Bay Area, ahead of several of its peers, such as Cohere, Hugging Face, Inflection and Perplexity.^[17]

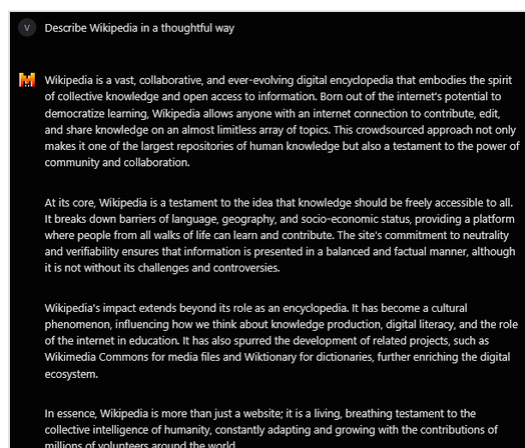
Partnership with Microsoft

On 26 February 2024, Microsoft announced a new partnership with the company to expand its presence in the artificial intelligence industry.

Under the agreement, Mistral's language models will be available on Microsoft's Azure cloud, while the multilingual conversational assistant *Le Chat* will be launched in the style of ChatGPT.^[18]



Example of an image generated with *Le Chat*, the prompt is Generate an image you feel represents yourself, Mistral AI



Screenshot of *Le Chat*, Mistral AI chatbot, describing Wikipedia in a thoughtful way

Services

Le Chat

On November 19, 2024, the company announced updates for *Le Chat*.

It added the ability to create images, in partnership with Black Forest Labs, utilizing the Flux Pro model.

Additionally, it introduced the capability to search for information on the internet to provide reliable and up-to-date information.

Furthermore, it launched the Canvas system, a collaborative interface where the AI generates code and the user can modify it.

Mobile app

On February 6, 2025, Mistral AI released its AI assistant, *Le Chat*, on iOS and Android, making its language models accessible on mobile devices.

Le Chat offers features including web search, image generation, and real-time updates.

Mistral AI also introduced a Pro subscription tier, priced at \$14.99 per month, which provides access to more advanced models, unlimited messaging, and web browsing.^[19]

Models

Name	Release Date	Number of Parameters (Billion)	License	Notes
Mistral Small 3 <i>25.01</i>	January 2025	24	Apache 2.0	Upon its release in January 2025, Mistral Small 3 is benchmarked as the leader in the "small" models category below 70B, featuring 24B parameters and capabilities comparable to those of larger models. ^{[20][21]}
Mistral Large 2 <i>24.11</i>	November 2024	123	Mistral Research License	^[21]
Pixtral Large <i>24.11</i>	November 2024	124	Mistral Research License	On November 19, 2024, the company introduced Pixtral Large, an improvement over Pixtral 12B, integrating a 1-billion-parameter visual encoder coupled with Mistral Large 2. This model has also been enhanced, particularly for long contexts and function calls. ^{[22] [21]}
Ministral 8B <i>24.10</i>	October 2024	8	Mistral Research License	^[21]
Ministral 3B <i>24.10</i>	October 2024	3	Proprietary	^[21]
Pixtral <i>24.09</i>	September 2024	12	Apache 2.0	^[21]
Mistral Large 2 <i>24.07</i>	July 2024	123	Mistral Research License	Mistral Large 2 was announced on July 24, 2024, and released on Hugging Face. It is available for free with a Mistral Research Licence, and with a commercial licence for commercial purposes. Mistral AI claims that it is fluent in dozens of languages, including many programming languages. Unlike the previous Mistral Large, this version was released with open weights. The model has 123 billion parameters and a context length of 128,000 tokens. ^[21]
Codestral Mamba 7B	July 2024	7	Apache 2.0	Codestral Mamba is based on the Mamba 2 architecture, which allows it to generate responses even with longer input. ^[23] Unlike Codestral, it was released under the Apache 2.0 license. While previous releases often included both the base model and the instruct version, only the instruct version of Codestral Mamba was released. ^{[24][21]}
Mathstral 7B	July 2024	7	Apache 2.0	Mathstral 7B is a model with 7 billion parameters released by Mistral AI on July 16, 2024, focusing on STEM subjects. ^[25] The model was produced in collaboration with Project Numina, ^[23] and was released under the Apache 2.0 License with a context length of 32k tokens. ^{[25] [21]}

Name	Release Date	Number of Parameters (Billion)	License	Notes
Codestral 22B	May 2024	22	Mistral Non-Production License	Codestral is Mistral's first code-focused open weight model which was launched on May 29, 2024. Mistral claims Codestral is fluent in more than 80 programming languages ^[26] Codestral has its own license which forbids the usage of Codestral for commercial purposes. ^{[27][21]}
Mistral 8x22B	April 2024	22	Apache 2.0	Similar to Mistral's previous open models, Mixtral 8x22B was released via a BitTorrent link on Twitter on April 10, 2024, ^[28] with a release on Hugging Face soon after. ^[29] The model uses an architecture similar to that of Mistral 8x7B, but with each expert having 22 billion parameters instead of 7. In total, the model contains 141 billion parameters, as some parameters are shared among the experts, but offering higher performance. ^{[29][30] [21]}
Mistral Small	February 2024	Unknown	Proprietary	Like the Large model, Mistral Small was launched on February 26, 2024. ^[21]
Mistral Large 24.02	February 2024	Unknown	Proprietary	Mistral Large was launched on February 26, 2024, and Mistral claims it is second in the world only to OpenAI's GPT-4. It is fluent in English, French, Spanish, German, and Italian, with Mistral claiming understanding of both grammar and cultural context, and provides coding capabilities. As of early 2024, it is Mistral's flagship AI. ^[31] It is also available on Microsoft Azure. ^{[32][21]}
Mistral Medium	December 2023	Unknown	Proprietary	Mistral Medium is trained in various languages including English, French, Italian, German, Spanish and code with a score of 8.6 on MT-Bench. ^[33] It is ranked in performance above Claude and below GPT-4 on the LMSys ELO Arena benchmark. ^[34] The number of parameters, and architecture of Mistral Medium is not known as Mistral has not published public information about it. ^[21]
Mistral 8x7B	December 2023	46.7	Apache 2.0	Much like Mistral's first model, Mixtral 8x7B was released via a BitTorrent link posted on Twitter on December 9, 2023, ^[2] and later Hugging Face and a blog post were released two days later. ^[35] Unlike the previous Mistral model, Mixtral 8x7B uses a sparse mixture of experts architecture. The model has 8 distinct groups of "experts", giving the model a total of 46.7B usable parameters. ^{[36][37]} Each single token can only use 12.9B parameters, therefore giving the speed and cost that a 12.9B parameter model would incur. ^[35] A version trained to follow instructions called "Mixtral 8x7B Instruct" is also offered. ^{[35][21]}
Mistral 7B	September 2023	7.3	Apache 2.0	Mistral 7B is a 7.3B parameter language model using the transformers architecture. It was officially released on September 27, 2023, via a BitTorrent magnet link, ^[38] and Hugging Face ^[39] under the Apache 2.0 license. Mistral 7B employs grouped-query attention (GQA), which is a variant of the standard attention mechanism. This architecture optimizes performance by calculating attention within specific groups of hidden states rather than across all hidden states, improving efficiency and scalability. ^[40] Both a base

Name	Release Date	Number of Parameters (Billion)	License	Notes
				model and "instruct" model were released with the latter receiving additional tuning to follow chat-style prompts. The fine-tuned model is only intended for demonstration purposes, and does not have guardrails or moderation built-in. ^{[41][21]}

Performance

Mistral 7B

Mistral AI claimed in the Mistral 7B release blog post that the model outperforms LLaMA 2 13B on all benchmarks tested, and is on par with LLaMA 34B on many benchmarks tested,^[41] despite having only 7 billion parameters, a small size compared to its competitors.

Mixtral 8x7B

Mistral AI's testing in 2023 shows the model beats both LLaMA 70B, and GPT-3.5 in most benchmarks.^[42]

In March 2024, a research conducted by Patronus AI comparing performance of LLMs on a 100-question test with prompts to generate text from books protected under U.S. copyright law found that Open AI's GPT-4, Mixtral, Meta AI's LLaMA-2, and Anthropic's Claude 2 generated copyrighted text verbatim in 44%, 22%, 10%, and 8% of responses respectively.^{[43][44]}

Mistral Large 2

According to Mistral AI, Large 2's performance in benchmarks is competitive with Llama 3.1 405B, particularly in programming-related tasks.^{[45][46]}

Codestral 22B

As of its release date, Codestral 22B surpasses Meta's Llama3 70B and DeepSeek Coder 33B (78.2% - 91.6%), another code-focused model on the HumanEval FIM benchmark.^[47]

Mathstral 7B

Mathstral 7B achieved a score of 56.6% on the MATH benchmark and 63.47% on the MMLU benchmark.^[25]

Usage

According to Mistral AI,^[48] the company's products have been used by :

1. BNP Paribas
2. AXA
3. Laboratoires Pierre Fabre
4. CMA CGM
5. Zalando
6. Mirakl
7. France Travail


References

1. Bradshaw, Tim; Abboud, Leila (30 January 2025). "Has Europe's great hope for AI missed its moment?" (<https://www.ft.com/content/fa8bad75-dc55-47d9-9eb4-79ac94e54d82>). *Financial Times*.
2. "Buzzy Startup Just Dumps AI Model That Beats GPT-3.5 Into a Torrent Link" (<https://gizmodo.com/mistral-artificial-intelligence-gpt-3-openai-1851091217>). Gizmodo. 12 December 2023. Retrieved 16 December 2023.
3. "What is Mistral AI?" (<https://www.ibm.com/think/topics/mistral-ai>). *IBM*. October 2024.
4. Journal, Sam Schechner | Photographs by Edouard Jacquinet for The Wall Street. "The 9-Month-Old AI Startup Challenging Silicon Valley's Giants" (<https://www.wsj.com/tech/ai/the-9-month-old-ai-startup-challenging-silicon-valleys-giants-ee2e4c48>). *WSJ*. Retrieved 31 March 2024.
5. "Spotlight Interview: Mistral AI CEO Arthur Mensch" (<https://frenchtechjournal.com/spotlight-interview-mistral-ai-arthur-mensch/>). *The French Tech Journal*. 24 June 2024.
6. "France's unicorn start-up Mistral AI embodies its artificial intelligence hopes" (https://www.lemonde.fr/en/economy/article/2023/12/12/french-unicorn-start-up-mistral-ai-embodies-its-artificial-intelligence-hopes_6337125_19.html). *Le Monde.fr*. 12 December 2023.
7. "Bringing open AI models to the frontier" (<https://mistral.ai/news/about-mistral-ai/>). Mistral AI. 27 September 2023. Retrieved 4 January 2024.
8. Webb, Maria (2 January 2024). "Mistral AI: Exploring Europe's Latest Tech Unicorn" (<https://www.techopedia.com/mistral-ai-exploring-europes-latest-tech-unicorn>). *techopedia.com*. Retrieved 13 June 2024.
9. "Mistral lève 385 M€ et devient une licorne française - le Monde Informatique" (<https://www.lemondeinformatique.fr/actualites/lire-mistral-leve-385-meteuro-et-devient-une-licorne-francaise-92392.html>). 11 December 2023.
10. Metz, Cade (10 December 2023). "Mistral, French A.I. Start-Up, Is Valued at \$2 Billion in Funding Round" (<https://www.nytimes.com/2023/12/10/technology/mistral-ai-funding.html>). *The New York Times*.
11. Fink, Charlie. "This Week In XR: Epic Triumphs Over Google, Mistral AI Raises \$415 Million, \$56.5 Million For Essential AI" (<https://www.forbes.com/sites/charlifink/2023/12/14/this-week-in-xr-epic-triumphs-over-google-mistral-ai-raises-415-million-565-million-for-essential-ai/>). *Forbes*. Retrieved 16 December 2023.
12. "A French AI start-up may have commenced an AI revolution, silently" (<https://www.hindustantimes.com/business/a-french-ai-start-up-may-have-commenced-an-ai-revolution-silently-101702370816617.html>). Hindustan Times. 12 December 2023.
13. Abboud, Leila; Levingston, Ivan; Hammond, George (8 December 2023). "French AI start-up Mistral secures €2bn valuation" (<https://www.ft.com/content/ea29ddf8-91cb-45e8-86a0-f501ab7ad9bb>). *Financial Times*. ft.com Financial Times.
14. Abboud, Leila; Levingston, Ivan; Hammond, George (19 April 2024). "Mistral in talks to raise €500mn at €5bn valuation" (<https://www.ft.com/content/358fcd8-8276-4e81-9ee6-e0391575d569>). *Financial Times*. Retrieved 19 April 2024.

15. Kharpal, Arjun (24 May 2024). "CEOs of AI startups backed by Microsoft and Amazon are the new tech rockstars" (<https://www.cnbc.com/2024/05/24/mistral-ai-anthropic-ceos-are-the-new-tech-rockstars-in-ai-boom.html>). *CNBC*. Retrieved 13 June 2024.
16. "Tripling Down on Mistral AI | General Catalyst" (<https://www.generalcatalyst.com/perspectives/tripling-down-on-mistral-ai>). *www.generalcatalyst.com*. Retrieved 13 June 2024.
17. Bratton, Laura (12 June 2024). "OpenAI's French rival Mistral AI is now worth \$6 billion. That's still a fraction of its top competitors" (<https://qz.com/openai-mistral-ai-funding-valuation-microsoft-1851535049>). *Quartz* (publication). Retrieved 13 June 2024.
18. Bableswhar (26 February 2024). "Mistral Large, Mistral AI's flagship LLM, debuts on Azure AI Models-as-a-Service" (<https://techcommunity.microsoft.com/t5/ai-machine-learning-blog/mistral-large-mistral-ai-s-flagship-llm-debuts-on-azure-ai/ba-p/4066996>). *techcommunity.microsoft.com*. Retrieved 26 February 2024.
19. "Mistral releases its genAI assistant Le Chat for IOS and Android" (<https://www.computerworld.com/article/3819932/mistral-releases-its-ai-assistant-le-chat-for-ios-and-android.html>). *Computerworld*. 7 February 2025. Retrieved 7 February 2025.
20. <https://ollama.com/library/mistral-small:24b>
21. https://docs.mistral.ai/getting-started/models/models_overview/}}
22. "Mistral has entered the chat" (<https://mistral.ai/news/mistral-chat/>). *Mistral AI*. 18 November 2024. Retrieved 11 December 2024.
23. David, Emilia (16 July 2024). "Mistral releases Codestral Mamba for faster, longer code generation" (<https://venturebeat.com/ai/mistral-releases-codestral-mamba-for-faster-longer-code-generation/>). *VentureBeat*. Retrieved 17 July 2024.
24. AI, Mistral (16 July 2024). "Codestral Mamba" (<https://mistral.ai/news/codestral-mamba/>). *mistral.ai*. Retrieved 16 July 2024.
25. AI, Mistral (16 July 2024). "MathΣtral" (<https://mistral.ai/news/mathstral/>). *mistral.ai*. Retrieved 16 July 2024.
26. Sharma, Shubham (29 May 2024). "Mistral announces Codestral, its first programming focused AI model" (<https://venturebeat.com/ai/mistral-announces-codestral-its-first-programming-focused-ai-model/>). *VentureBeat*. Retrieved 30 May 2024.
27. Wiggers, Kyle (29 May 2024). "Mistral releases Codestral, its first generative AI model for code" (<https://techcrunch.com/2024/05/29/mistral-releases-its-first-generative-ai-model-for-code/>). *TechCrunch*. Retrieved 30 May 2024.
28. @MistralAI (10 April 2024). "Torrent" (<https://x.com/MistralAI/status/1777869263778291896>) (Tweet) - via Twitter.
29. "mistralai/Mixtral-8x22B-v0.1 · Hugging Face" (<https://huggingface.co/mistralai/Mixtral-8x22B-v0.1>). *huggingface.co*. Retrieved 5 May 2024.
30. "Mistral Releases Latest Open Source Model, Mixtral 8x22B" (<https://pureai.com/Articles/2024/04/17/Mistral-Mixtral-8x22B.aspx>). *Pure AI*. 17 April 2024.
31. AI, Mistral (26 February 2024). "Au Large" (<https://mistral.ai/news/mistral-large/>). *mistral.ai*. Retrieved 6 March 2024.
32. Boyd, Eric (26 February 2024). "Introducing Mistral-Large on Azure in partnership with Mistral AI" (<https://azure.microsoft.com/en-us/blog/microsoft-and-mistral-ai-announce-new-partnership-to-accelerate-ai-innovation-and-introduce-mistral-large-first-on-azure/>). *Microsoft Azure Blog*. Retrieved 17 February 2025.
33. AI, Mistral (11 December 2023). "La plateforme" (<https://mistral.ai/news/la-plateforme/>). *mistral.ai*. Retrieved 22 January 2024.
34. "LMSys Chatbot Arena Leaderboard - a Hugging Face Space by lmsys" (<https://huggingface.co/spaces/lmsys/chatbot-arena-leaderboard>). *huggingface.co*. Retrieved 22 January 2024.
35. "Mixtral of experts" (<https://mistral.ai/news/mixtral-of-experts/>). *mistral.ai*. 11 December 2023. Retrieved 4 January 2024.

36. "Mixture of Experts Explained" (<https://huggingface.co/blog/moe>). *huggingface.co*. Retrieved 4 January 2024.
37. Marie, Benjamin (15 December 2023). "Mixtral-8x7B: Understanding and Running the Sparse Mixture of Experts" (<https://towardsdatascience.com/mixtral-8x7b-understanding-and-running-the-sparse-mixture-of-experts-0e3fc7fde818>). *Medium*. Retrieved 4 January 2024.
38. Goldman, Sharon (8 December 2023). "Mistral AI bucks release trend by dropping torrent link to new open source LLM" (<https://venturebeat.com/ai/mistral-ai-bucks-release-trend-by-dropping-torrent-link-to-new-open-source-llm/>). *VentureBeat*. Retrieved 4 January 2024.
39. Coldewey, Devin (27 September 2023). "Mistral AI makes its first large language model free for everyone" (<https://techcrunch.com/2023/09/27/mistral-ai-makes-its-first-large-language-model-free-for-everyone/>). *TechCrunch*. Retrieved 4 January 2024.
40. Jiang, Albert Q.; Sablayrolles, Alexandre; Mensch, Arthur; Bamford, Chris; Chaplot, Devendra Singh; Casas, Diego de las; Bressand, Florian; Lengyel, Gianna; Lample, Guillaume (10 October 2023). "Mistral 7B". *arXiv:2310.06825v1* (<https://arxiv.org/abs/2310.06825v1>) [cs.CL (<https://arxiv.org/archive/cs/CL>)].
41. "Mistral 7B" (<https://mistral.ai/news/announcing-mistral-7b/>). *mistral.ai*. Mistral AI. 27 September 2023. Retrieved 4 January 2024.
42. Franzen, Carl (11 December 2023). "Mistral shocks AI community as latest open source model eclipses GPT-3.5 performance" (<https://venturebeat.com/ai/mistral-shocks-ai-community-as-latest-open-source-model-eclipses-gpt-3-5-performance/>). *VentureBeat*. Retrieved 4 January 2024.
43. Field, Hayden (6 March 2024). "Researchers tested leading AI models for copyright infringement using popular books, and GPT-4 performed worst" (<https://www.cnn.com/2024/03/06/gpt-4-researchers-tested-leading-ai-models-for-copyright-infringement.html>). *CNN*. Retrieved 6 March 2024.
44. "Introducing CopyrightCatcher, the first Copyright Detection API for LLMs" (<https://www.patronus.ai/blog/introducing-copyright-catcher>). Patronus AI. 6 March 2024. Retrieved 6 March 2024.
45. AI, Mistral (24 July 2024). "Large Enough" (<https://mistral.ai/news/mistral-large-2407/>). *mistral.ai*. Retrieved 24 July 2024.
46. "mistralai/Mistral-Large-Instruct-2407 · Hugging Face" (<https://huggingface.co/mistralai/Mistral-Large-Instruct-2407>). *huggingface.co*. Retrieved 24 August 2024.
47. AI, Mistral (29 May 2024). "Codestral: Hello, World!" (<https://mistral.ai/news/codestral/>). *mistral.ai*. Retrieved 30 May 2024.
48. "Solutions - for any use case | Mistral AI" (<https://mistral.ai/en/solutions>). *mistral.ai*. Retrieved 17 February 2025.

External links

- [Official website](https://mistral.ai/) (<https://mistral.ai/>) 
 - [Official Fan Twitter](https://x.com/LeChatonSol) (<https://x.com/LeChatonSol>)
-

Retrieved from "https://en.wikipedia.org/w/index.php?title=Mistral_AI&oldid=1277140258"